


Deferred Rendering

Introduction to screen-space lighting

Programming – Computer Graphics

Contents

- The Problem of Forward Rendering
- Deferred Rendering
 - Steps
 - G-Pass
 - Light Pass
 - Composite Pass
 - Benefits
 - Optimisations



The Problem of Forward Rendering

- Forward Rendering is when we draw a mesh to the back-buffer, then the next mesh, and the next, each separate from the other
 - Pixels can be written to multiple times (mesh B renders over the pixels of mesh A, etc)
 - Overdraw
- This means our complex lighting calculations and shaders get wasted
 - Most PC/Console games are Fragment Shader heavy, rather than vertex heavy, so wasting calculations can be costly!
- We also only have access to the pixel currently being rendered
 - We'd have to use Post-Processing to sample nearby fragments, but this would be after lighting has been calculated


Deferred Rendering

- A popular method of rendering that exists in almost all major PC/Console engines is **Deferred Rendering**
 - Also called **Deferred Shading**
- A scene is rendered to off-screen buffers without applying lighting/shading, then the lighting is applied as a **Post-Processing** step
 - This reduces the lighting calculations by not having to perform them for pixels that would get written over with overdraw
 - Only visible pixels are lit



Deferred Rendering Steps

- There are 3 stages to deferred rendering
 - The Geometry Pass or G-Pass
 - The Light Pass
 - The Composite Pass



The G-Pass

- This stage is done as a forward pass
 - Each mesh is rendered into off-screen buffers rather than the Back Buffer, called the G-Buffers
- Traditionally renders into buffers that hold
 - Material information, such as unlit diffuse colour of geometry, called the **Albedo**
 - **Positional** information of the geometry at each pixel, or depth as distance from the camera
 - **Normal** information of the geometry at each pixel
- Requires the use of Multiple Render Targets



The G-Pass

- Once we have rendered the entire scene in to the G-Buffers we can use them as input textures to the **Light Pass** and **Composite Pass**
- We can store various data within the G-Buffers but there are a few considerations
 - How many Components / Elements?
 - Data type and size
- Some examples:
 - Using 3 buffers with RGB 32F format (3x 32-bit float values)

Albedo
Red
Green
Blue

Position
X
Y
Z

Normal
X
Y
Z


The Light Pass

- The next pass in Deferred Rendering is the Light Pass
 - We render in to a new render target, the **Light Buffer**
 - We use 2 of the render targets from the previous pass as textures; the **Position Buffer** and **Normal Buffer**
- We then render all lights as geometry, with additive blending enabled
 - We use geometry that would cover the area of the screen that the light would illuminate
 - Directional lights can be rendered as a full-screen quad
 - Point lights can be rendered as spheres or boxes
 - Spot lights can be rendered as cones or pyramids
 - The geometry is rendered with a standard Vertex Shader
 - The Fragments that get processed are then fragments which may be illuminated by the light




The Light Pass

- Within the Fragment Shader for the light geometry we calculate the light for the pixel
 - At each pixel we know its Position and Normal thanks to the G-Pass
 - We simply sample the textures for that pixel
 - We then calculate lighting as usual
 - The Normal comes from the Normal Buffer
 - Position from the Position Buffer can be used for calculating specular and view vectors
- We don't use the surface colour in the calculation, only the light colour
- We then output the colour with additive blending
 - The more lights shining on a single pixel the brighter it will be
- We end up with a buffer that contains all light for each pixel



The Composite Pass

- Once all lights have been rendered we can move on to the **Composite Pass**
 - This is the final pass and renders typically in to the Back Buffer
 - We use the **Albedo Buffer** and the **Light Buffer** from the earlier passes
- This pass is a simple Post-Process pass
 - We render a full-screen quad with the Albedo and Light Buffers applied as textures
 - The Fragment Shader simply multiplies the Albedo with the Light Buffer



Deferred Rendering Benefits

- The main benefit of Deferred Rendering is the reduction in lighting complexity and ability to have many lights
 - To be able to light a mesh with many lights in Forward Rendering we would have to render the mesh multiple times, or have shaders available for all the different light counts

$$O(\text{meshCount} \times \text{lightCount})$$

- With Deferred we can now have many lights as they simply render as geometry
 - Each mesh is rendered, then each light

$$O(\text{meshCount} + \text{lightCount})$$



Deferred Rendering Benefits

- This means we can have environments with hundreds or thousands of lights!
 - We also only need one set of lighting shaders, one for each type of light, rather than different shaders depending on the number of active lights in a scene



Deferred Rendering Benefits



Deferred Rendering Pitfalls

- Transparency is extremely hard to include!
 - Typically transparent objects are rendered afterwards in a forward rendering manner, back to front
 - Modern techniques allow forms of “layers” or pixel linked-lists in deferred rendering for transparent objects, but these can be slow and require high-end hardware
- Memory Footprint Issues
 - Not all systems like Multiple Render Targets
 - High Definition displays require large G-Buffers
 - Around 8mb for a single buffer using 4 bytes, closer to 32mb for a vec4-based buffer!
 - Video card memory can disappear fast
- All meshes use the same material!




Deferred Rendering Optimisations

- We can reduce the size of the data stored in the g-buffer
 - Normals usually take 3 elements, XYZ, and are best stored as floats for accuracy
 - Position also would typically use 3 elements for XYZ
- Rather than position we could store depth as a single float
 - Using depth we could recreate the position by projecting the 2D screen-coordinate back into 3D at the required depth
- Rather than XYZ for the normal we could store 2 floats!
 - Early implementations ignored Z with the assumption that in view-space it always faced the camera and so was positive, and could then be recreated
 - However this is incorrect!
 - Newer implementations convert the normal to other formats, such as 2 polar angles, or to sphere-map coordinates with Azimuthal Projections



Summary

- Deferred Rendering allows for many lights in a scene
 - Thousands!
- Performance issues must be taken into account, but optimisations can reduce them substantially
- Majority of AAA developers use deferred techniques, as do majority of commercial game engines
 - Unreal, Unity, Source, Frostbite, PhyreEngine, CryEngine



Further Reading

- Wolff, D, 2013, *OpenGL 4 Shading Language Cookbook*, 2nd Edition, Chapter 5, PACKT Publishing
- Haemel, N, Sellers, G & Wright, R, 2014, *OpenGL SuperBible*, 6th Edition, Chapter 12, Addison Wesley