


Introduction to the Graphics Processing Unit

An overview of Computer Graphics, Render Pipelines and APIs

Programming – Computer Graphics

Computer Graphics

- Computer graphics is any sort of visualisation created and displayed by a computer
- Early imagery from the 1950's created vector-based displays using series of lines
- Methods of rasterising bitmap imagery and geometry enabled images with colour and texture variation
- Eventually dedicated hardware was created to calculate the costly methods needed to display complex graphics
 - This created an explosion in visual fidelity



The Graphics Processing Unit (GPU)

- Initially developed to increase visual performance when drawing basic shapes and speed up bitmap manipulations
- Processors capable of 3D manipulation and Hardware Texture & Lighting (T&L) were eventually released
- Each manufacturer implemented their own way for rendering graphics
 - This caused problems for PC game developers dealing with the vast array of hardware configurations



Introduction of Standards

- The need for standards was agreed upon and a set render pipeline was formalised with API developers leading the charge
- Multiple APIs capable of manipulating these units
 - OpenGL
 - Direct3D (part of the DirectX set of APIs)
 - Vulkan
 - Metal




APIs

- The various APIs are a way for programmers to interface with the GPU drivers without having to write code specific for each device driver
 - The API developers and the hardware manufacturers make sure they work for us
- APIs do have limitations though
 - Direct3D only applies to Microsoft operating systems
 - OpenGL is cross-platform across PC hardware
 - OpenGL ES is commonly used on mobile devices
 - Metal only applies to Apple operating systems

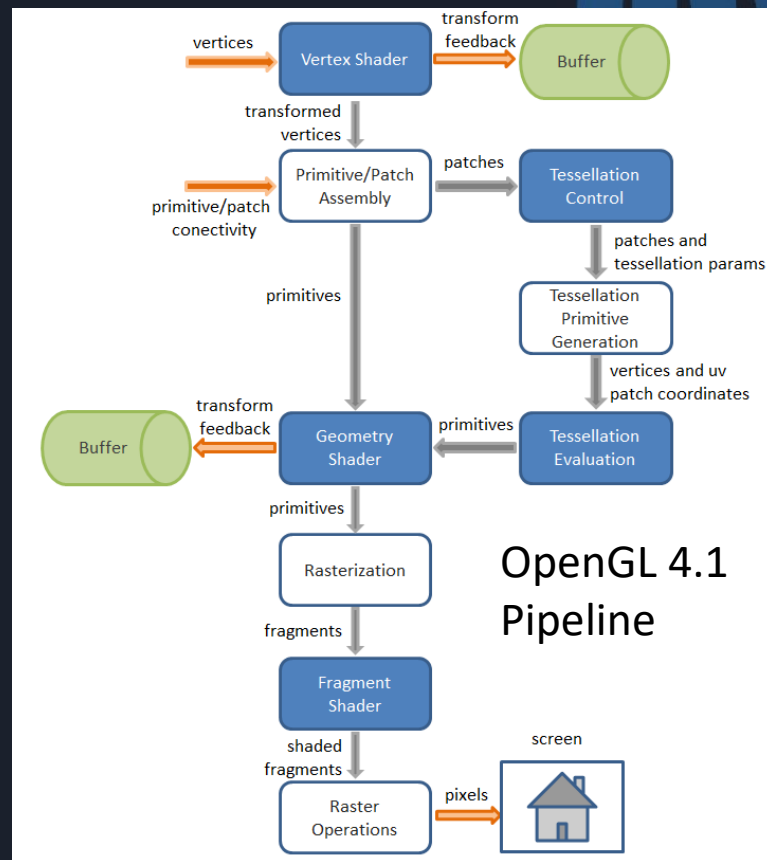
The Early Render Pipeline

- A render pipeline refers to the method and stages of rasterisation for graphics hardware
 - Rasterisation refers to converting vector-based graphic images into a raster image consisting of pixels
 - Geometric data is input into the pipeline, and coloured pixels are drawn as output
- In the beginning the amount of customisation was limited
 - An artist could create a mesh and textures, choose colours for a light or two, but that was all that could be done to differentiate the look of their game from other games



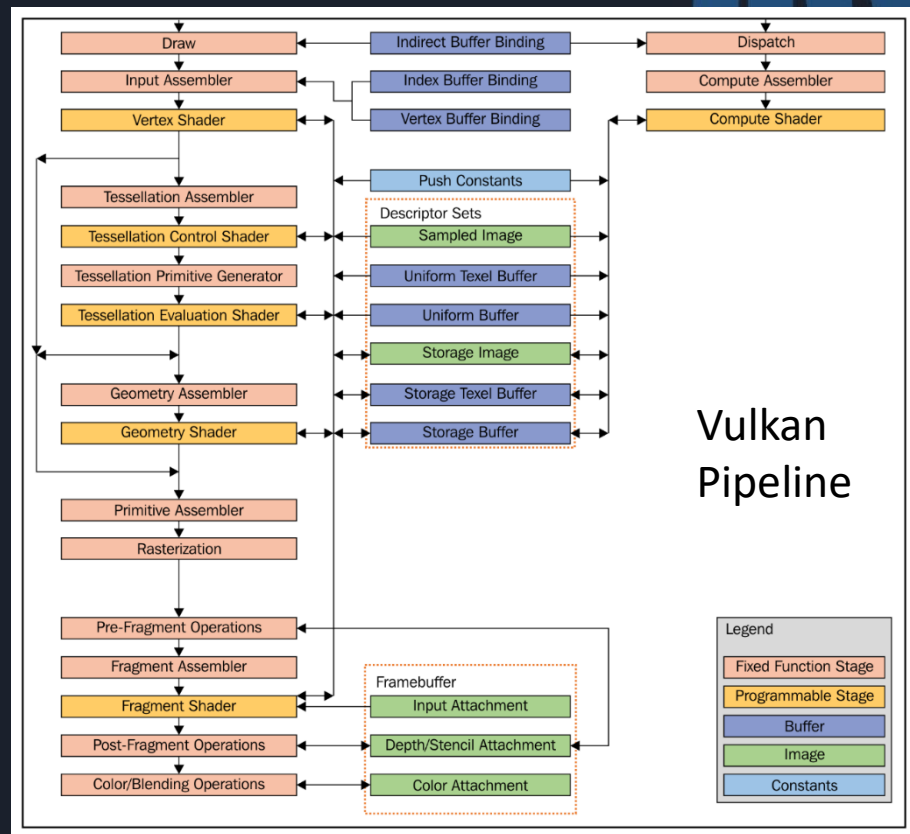
2000-2012 Render Pipeline

- GPU hardware improved greatly, with new features allowing for more complex render pipelines
 - Increased GPU speed
 - Increased GPU RAM
 - Increased register and operation counts
 - Programmable and Fixed stages
- Game developers could now completely customise how their products look



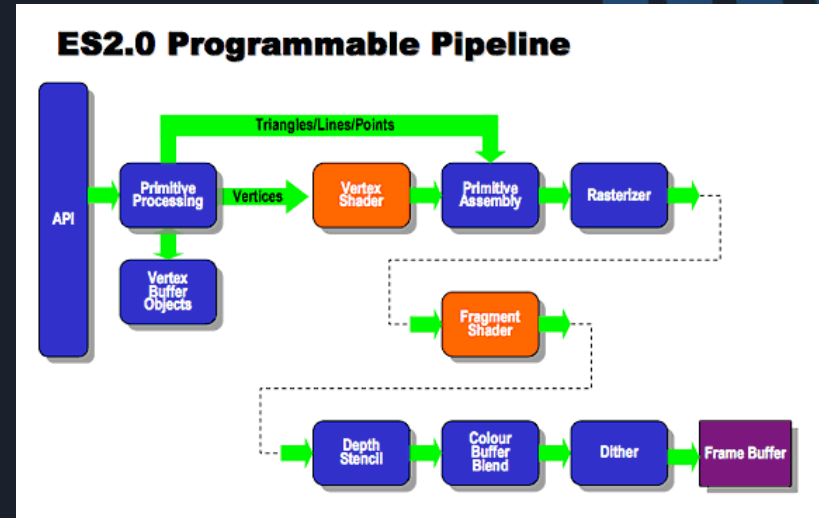
Modern Render Pipelines

- Current generation pipelines are complex with many programmable stages
- Data can now flow both ways between the GPU and CPU at multiple points in the pipeline via buffers
 - Many inputs and many outputs, not just input geometry and output pixels
- Compute stages allow custom rendering or advanced computation of almost any type of data



Offshoot Render Pipelines

- Other pipelines exist
 - OpenGL ES (Embedded Systems)
 - WebGL for web sites (resembles OpenGL ES)
- These are usually cut-down versions of standard pipelines and contain less stages due to hardware limitations on target devices




Buffer Objects

- The inputs and outputs of a typical Render Pipeline are called **Buffer Objects**
- There are different types of buffers, usually they are simply arrays of data
 - **Vertex Buffer Object**, containing the points that make up geometry for a mesh
 - **Index Buffer Object**, containing topology information for the geometry, i.e. triangles
 - Various image buffers, usually containing image pixel data, which can include data representing the screen being displayed or textures mapped onto geometry
- Buffers can be used as input to various stages
 - Vertex and Index buffers are used as input to the pipeline
 - Image buffers can be accessed at different points through the pipeline
- Some stages in a render pipeline may also output buffers
 - For example, the final output of the pipeline usually goes into an image buffer

The Vertex Shader Stage

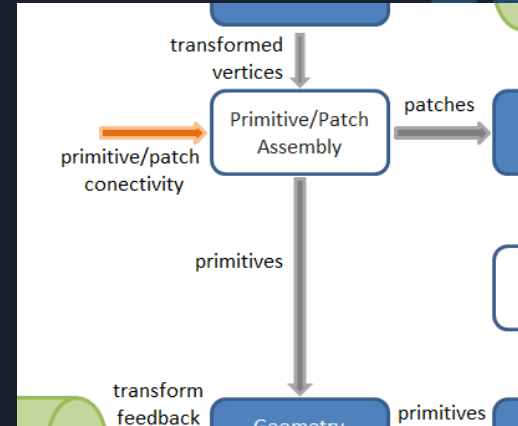
- Once the buffers have been setup and sent to the GPU the first stage they usually enter when we make a render call is the **Vertex Shader** stage



- This is a programmable stage that processes each vertex in a vertex buffer individually and separate from each other
 - Multiple vertices are processed at once through simple functions
- Typically the task of the Vertex Shader is to transform vertices from object space to screen space for later stages in the pipeline

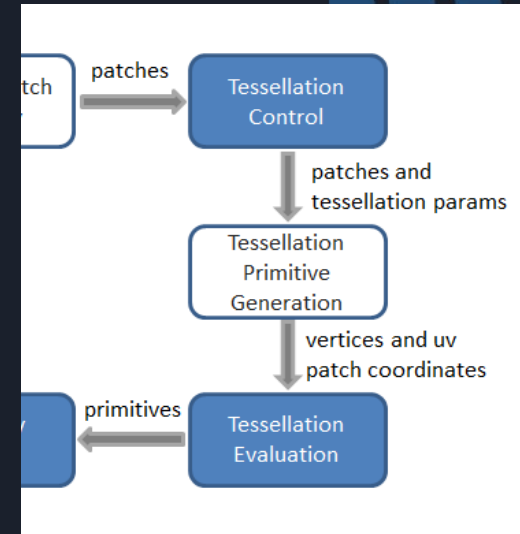
Primitive / Patch Assembly Stage

- The next stage is a fixed stage, and it is responsible for arranging the vertex data into primitives for the later stages
 - Makes use of vertex and index buffers
- This stage is usually controlled simply by specifying the primitive type when rendering
 - Triangles, lines, points, patches
- Outputs to either an optional Tessellation stage or to the Geometry Shader stage



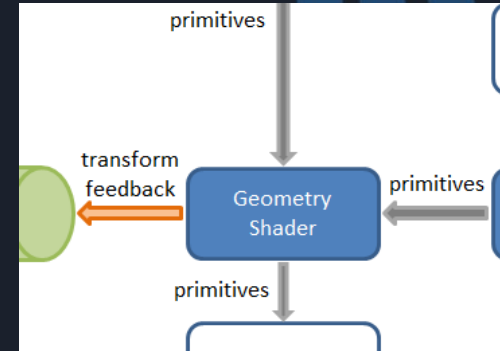
Tessellation Stages

- These are an optional set of three stages
 - **Tessellation Control** Stage (Programmable)
 - **Tessellation Primitive Generation** Stage (Fixed)
 - **Tessellation Evaluation** Stage (Programmable)
- The tessellator takes in patch primitives and splits them into multiple primitives, increasing the primitive count in a mesh
 - A patch primitive can have 1 or more vertices (32-max usually)
 - The **Control** stage defines how many times to split up the input patch
 - The **Generation** stage then generates the new primitives
 - The **Evaluation** stage receives these new primitives and then sends them through to the next stage in the render pipeline



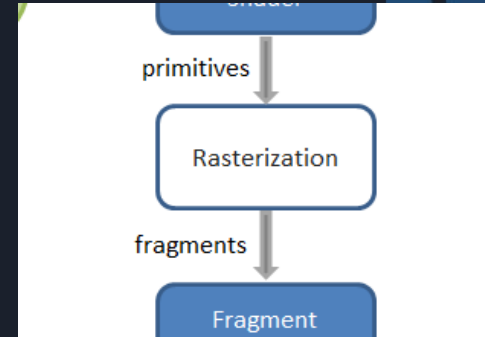



Geometry Shader Stage

- An optional programmable stage
 - If not set then it passes primitives straight to the Rasterisation stage
- Receives all the information needed for a primitive
 - For example, 3 vertices in the case of a triangle, 2 for a line
- Can perform any last minute processing on the primitive




Rasterisation & Interpolation Stage

- This fixed stage receives vector-based primitives and plots out the pixels covered by the shape




- It then interpolates the data from each vertex across the pixels covered, passing the interpolated data at each pixel to the next stage in the pipeline




Fragment Shader Stage

- This is the final programmable stage in a typical Render Pipeline, but not the last stage
- Also called the **Pixel Shader**, as it receives the interpolated pixel data from the rasteriser
- Typically its job is to output a desired colour at the specified pixel, based on the input data, and by sampling other buffers usually containing texture information
 - Can output more than one pixel to more than one output buffer!
- However the data returned is not necessarily what appears on screen!



Raster Operations Stage

- The final stage in the pipeline is a fixed stage
- Receives data from the **Fragment Shader**
- Using flags, it blends the pixel with any existing pixel at that location within the output buffer



Final Output Data

- The render pipeline's job is to output all the pixels represented by geometry
 - The pixels of all geometry in a level can be combined into a final image buffer
- Typically the output buffer represents the screen but it can also output to other locations
 - Data can be rendered into an image to be used as a texture when rendering other geometry



Summary

- The render pipeline has many parts to it, and people have found numerous different ways of using it to achieve visuals of all types
 - Some of its most recent uses aren't even for visuals!
- The best way to understand it is to just tinker with it and see what types of visuals and effects you can come up with!

Further Reading

- Akenine-Möller, T, Haines, E & Hoffman, N, 2008, *Real-Time Rendering*, 3rd Ed, CRC Press
- Haemel, N, Sellers, G & Wright, R, 2014, *OpenGL SuperBible*, 6th Ed, Addison Wesley